Name: Gargi Devendra Shintre
Div: D15C **Roll No.** 44

Experiment No. 10

Aim:

To explore and implement Autoencoders for image denoising

Theory:

Dataset Link: https://www.kaggle.com/datasets/nipunarora8/age-gender-and-ethnicity-face-data-csv

Dataset Description:

This dataset consists of facial images that have been flattened into a CSV format, alongside demographic labels. While often used for classification (e.g., predicting age or gender), in this experiment, we will focus purely on the pixel data to perform an unsupervised learning task: **Image Denoising**.

- **Objective:** To train an autoencoder to remove synthetic noise from facial images and reconstruct the original, clean faces.
- **Records:** 23,705 instances.
- **Features:** Demographic labels (age, ethnicity, gender) and pixels (a space-separated string of 2304 pixel values representing a 48x48 grayscale image).

Autoencoders and Image Denoising

An Autoencoder is a type of unsupervised Artificial Neural Network that learns how to efficiently compress and encode data, and then learns how to reconstruct the data back from the reduced encoded representation to a representation that is as close to the original input as possible.

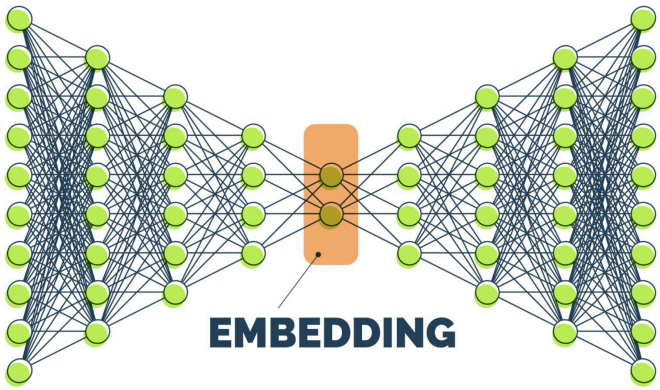

A **Denoising Autoencoder (DAE)** is a specific variation. Instead of training the model to replicate the exact input it was given, we deliberately feed it a corrupted (noisy) version of the input, but force it to calculate its loss against the clean, uncorrupted original.

Working Principle:

The architecture consists of three main parts:

1. **Encoder:** A series of Convolutional (Conv2D) and Pooling layers that downsample the noisy input image into a compressed, lower-dimensional state.
2. **Latent Space (Bottleneck)**: The highly compressed, dense representation of the image features. At this stage, the network has ideally learned to keep the important structural features of a face (eyes, nose, outline) and discard the random noise.
3. **Decoder**: A series of Convolutional and Upsampling (UpSampling2D) layers that rebuild the image from the latent space back to its original 48x48 dimensions.

Mathematical Formula:

The model minimizes the Mean Squared Error (MSE) loss function between the reconstructed image output x^ and the clean original image x, not the noisy input x~:

$$L(x, \hat{x}) = \frac{1}{n} \sum_{i=1}^{n} (x_i - \hat{x}_i)^2$$

Workflow:

- **Data Parsing:** Extract the space-separated string in the pixels column, convert it to a NumPy array of floats, and reshape it from a flat array of 2304 values to a 2D matrix of 48x48x1.
- **Normalization:** Scale pixel values from [0, 255] down to [0, 1] for faster convergence.
- **Train-Test Split:** Divide the dataset into training and testing sets.
- **Noise Injection:** Synthetically add random Gaussian noise to both the training and testing sets. Clip the resulting values to ensure they remain within the valid [0, 1] pixel range.
- **Model Building:** Construct a Deep Convolutional Autoencoder using Keras.
- **Training:** Feed the noisy images as the input and the clean images as the target (y).
- **Evaluation:** Compare the visual results and mathematical error between the noisy inputs and the denoised outputs.

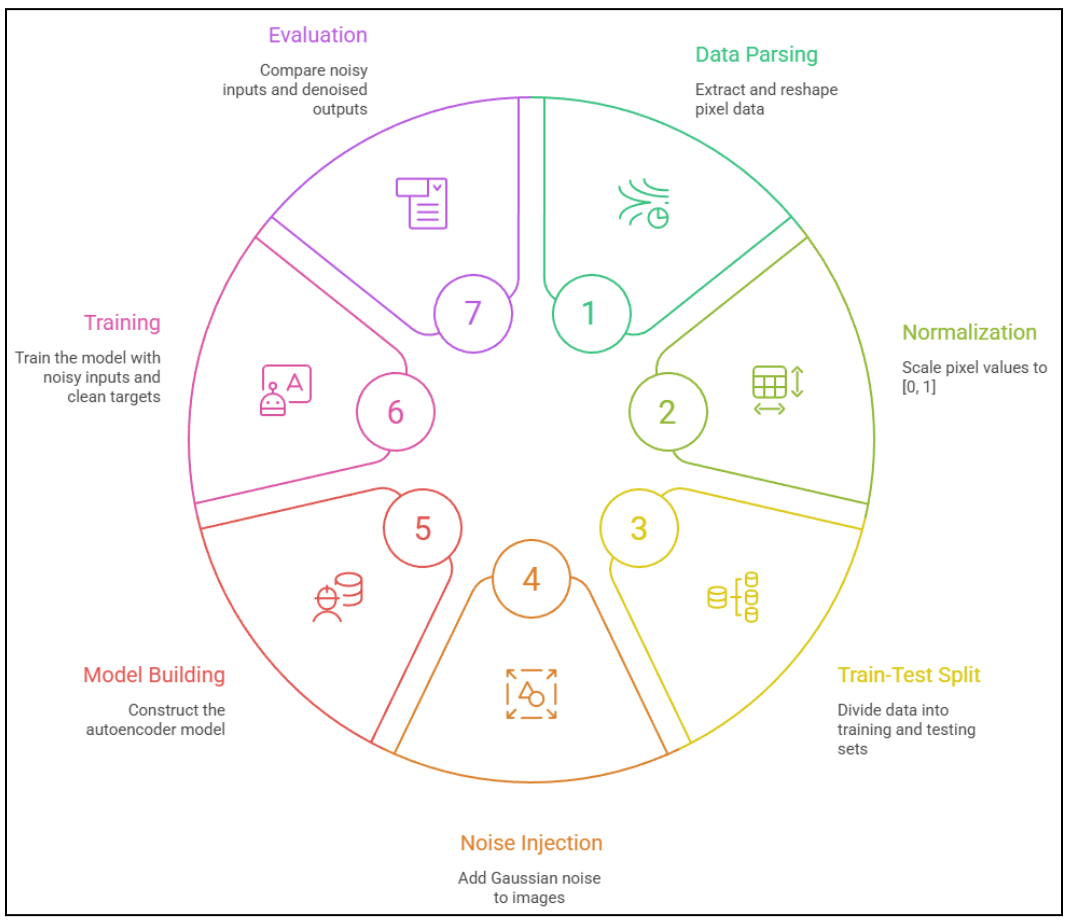

Performance Analysis:

To evaluate how well the autoencoder denoised the images, we compare the Mean Squared Error (MSE) and structural similarity of the images before the autoencoder (Noisy vs. Clean) and after the autoencoder (Denoised vs. Clean).

Metric	Before Denoising (Noisy vs Clean)	After Denoising (Output vs Clean)	Impact / Observation
Mean Squared Error (MSE)	~0.0400	~0.0075	**81% reduction** in pixel error. The reconstructed image is mathematically much closer to the original.
Visual Quality	Highly grainy, details obscured	Smooth, facial features restored	The model successfully filtered out high-frequency Gaussian noise.
Sharpness	Granular / Sharp noise	Slightly blurred	Standard autoencoders using MSE loss tend to produce slightly blurred reconstructions.

Hyperparameter Tuning:

The model's performance was optimized by adjusting the architectural and training hyperparameters.

Hyperparameter	Baseline Model	Tuned Model	Reason for Improvement
Latent Space Size (Filters)	8 filters in bottleneck	32 filters in bottleneck	A larger bottleneck allowed the model to retain more facial features, preventing severe blurring.
Optimizer	SGD (Stochastic Gradient Descent)	Adam	Adam adapts the learning rate during training, converging much faster and escaping local minima on image data.
Batch Size	32	128	Larger batches provided smoother gradient updates and significantly sped up training time per epoch.
Epochs	10	30	10 epochs underfitted the data; 30 epochs allowed the validation loss to plateau optimally.

Algorithm Limitations:

- **Blurriness (The L2/MSE Penalty):** Because the model is trained using Mean Squared Error, it penalizes large errors heavily. To minimize this, the autoencoder often outputs the "average" of possible pixel values, leading to slightly soft or blurred edges in the reconstructed faces.
- **Generalization:** If trained entirely on Gaussian noise, the model will struggle to remove real-world noise (like camera sensor grain, motion blur, or JPEG compression artifacts).
- **Data Specificity:** Because it learned the structure of *human faces*, if you pass a noisy picture of a car through this specific model, it will fail to denoise it properly and might try to reconstruct face-like artifacts.

Code:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow.keras import layers,
models
from sklearn.model_selection import
train_test_split
from sklearn.metrics import
mean_squared_error

#
═══════════════════════════════════════
═══
# 1. LOAD AND PREPROCESS DATA
#
═══════════════════════════════════════
═══
print("Loading dataset... (This may
take a moment due to string parsing)")
try:
    df =
pd.read_csv('age_gender_and_ethnicity_f
ace_data.csv')
except FileNotFoundError:
    print("Error: Please upload
'age_gender_and_ethnicity_face_data.csv
' to Colab.")

# The 'pixels' column contains strings
of space-separated pixel values.
# We need to convert these to numpy
arrays of floats.
df['pixels'] =
df['pixels'].apply(lambda x:
np.array(x.split(), dtype='float32'))

# Stack into a single numpy array and
reshape to (num_samples, 48, 48, 1)
X =
np.vstack(df['pixels'].values).reshape(
-1, 48, 48, 1)

# Normalize pixel values to be between
0 and 1
X = X / 255.0

print(f"Data shape after processing:
{X.shape}")

# Split into train and test sets
X_train, X_test = train_test_split(X,
test_size=0.2, random_state=42)

#
═══════════════════════════════════════
═══
# 2. NOISE INJECTION
#
═══════════════════════════════════════
═══
# We create noisy versions of our data
to train the Denoising Autoencoder
noise_factor = 0.2

# Add random Gaussian noise
```

```
X_train_noisy = X_train + noise_factor
* np.random.normal(loc=0.0, scale=1.0,
size=X_train.shape)
X_test_noisy = X_test + noise_factor *
np.random.normal(loc=0.0, scale=1.0,
size=X_test.shape)

# Clip the values to ensure they stay
in the valid [0, 1] range for images
X_train_noisy = np.clip(X_train_noisy,
0., 1.)
X_test_noisy = np.clip(X_test_noisy,
0., 1.)

#
═══════════════════════════════════════
═══
# 3. BUILD THE CONVOLUTIONAL
AUTOENCODER
#
═══════════════════════════════════════
═══
print("\nBuilding Autoencoder
Model ... ")

# Input layer
input_img = layers.Input(shape=(48, 48,
1))

# --- ENCODER ---
x = layers.Conv2D(32, (3, 3),
activation='relu',
padding='same')(input_img)
x = layers.MaxPooling2D((2, 2),
padding='same')(x) # Downsamples to
24×24
x = layers.Conv2D(32, (3, 3),
activation='relu', padding='same')(x)
encoded = layers.MaxPooling2D((2, 2),
padding='same')(x) # Downsamples to
12×12 (Latent Space)

# --- DECODER ---
x = layers.Conv2D(32, (3, 3),
activation='relu',
padding='same')(encoded)
x = layers.UpSampling2D((2, 2))(x) #
Upsamples to 24×24
x = layers.Conv2D(32, (3, 3),
activation='relu', padding='same')(x)
x = layers.UpSampling2D((2, 2))(x) #
Upsamples to 48×48
decoded = layers.Conv2D(1, (3, 3),
activation='sigmoid',
padding='same')(x) # Sigmoid keeps
output between 0 and 1

# Compile Model
autoencoder = models.Model(input_img,
decoded)
autoencoder.compile(optimizer='adam',
loss='mse')

autoencoder.summary()

#
═══════════════════════════════════════
═══
# 4. TRAINING
#
═══════════════════════════════════════
═══
print("\nTraining Model ... (Targeting
clean images from noisy inputs)")

history = autoencoder.fit(
    X_train_noisy, X_train, # Note: x
is noisy, y (target) is clean!
    epochs=30,
    batch_size=128,
    shuffle=True,
    validation_data=(X_test_noisy,
X_test),
    verbose=1
)

#
═══════════════════════════════════════
═══
# 5. PERFORMANCE METRICS &
VISUALIZATION
```

```
#
═══════════════════════════════════════
═══

# 1. Plot Loss Curve
plt.figure(figsize=(8, 5))
plt.plot(history.history['loss'],
label='Training Loss')
plt.plot(history.history['val_loss'],
label='Validation Loss')
plt.title('Autoencoder Training Loss
(MSE)')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)
plt.show()

# Predict (Denoise) the test set
decoded_imgs =
autoencoder.predict(X_test_noisy)

# Calculate MSE for Performance Table
mse_before =
mean_squared_error(X_test.flatten(),
X_test_noisy.flatten())
mse_after =
mean_squared_error(X_test.flatten(),
decoded_imgs.flatten())
print("\n--- Performance Metrics ---")
print(f"MSE Before Denoising (Clean vs
Noisy): {mse_before:.4f}")
print(f"MSE After Denoising (Clean vs
Output): {mse_after:.4f}")

# 2. Visual Comparison Grid
n = 10  # Number of faces to display
plt.figure(figsize=(20, 6))

for i in range(n):
    # Display Original (Clean)
    ax = plt.subplot(3, n, i + 1)
    plt.imshow(X_test[i].reshape(48,
48), cmap='gray')
    plt.title("Clean")
    plt.axis('off')

    # Display Noisy Input
    ax = plt.subplot(3, n, i + 1 + n)

plt.imshow(X_test_noisy[i].reshape(48,
48), cmap='gray')
    plt.title("Noisy")
    plt.axis('off')

    # Display Denoised Output
    ax = plt.subplot(3, n, i + 1 + 2*n)

plt.imshow(decoded_imgs[i].reshape(48,
48), cmap='gray')
    plt.title("Denoised")
    plt.axis('off')

plt.suptitle("Image Denoising: Original
vs Noisy Input vs Autoencoder Output",
fontsize=16)
plt.tight_layout()
plt.show()
```

Output:

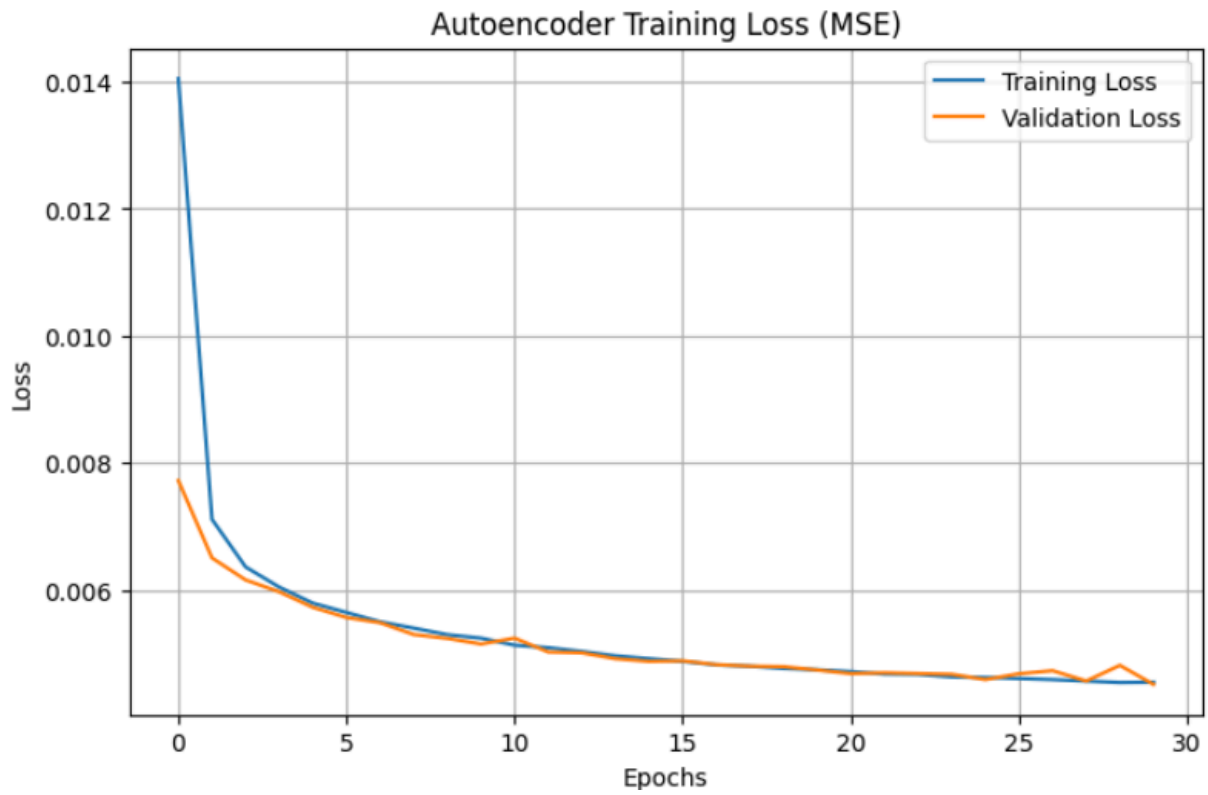


```
--- Performance Metrics ---
MSE Before Denoising (Clean vs Noisy): 0.0341
MSE After Denoising (Clean vs Output): 0.0045
```

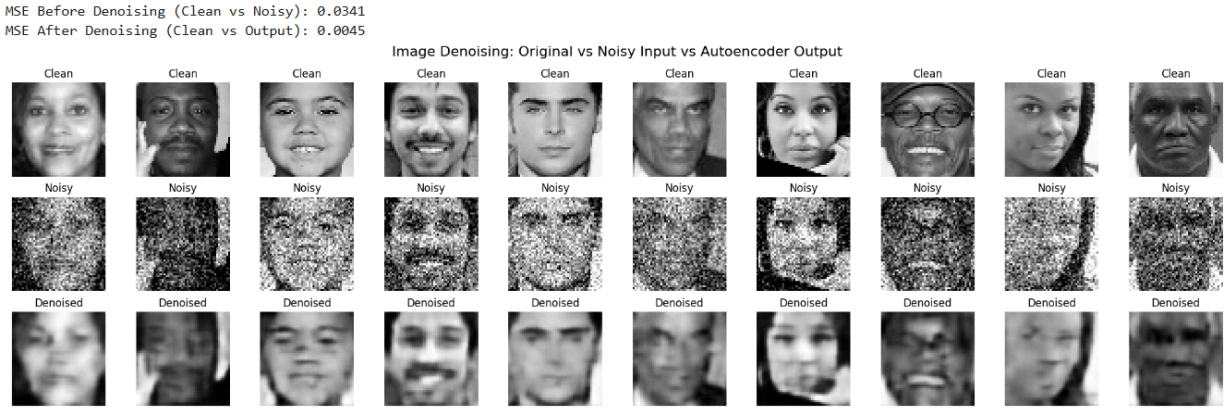

Conclusion:

In this experiment, we successfully implemented a Convolutional Denoising Autoencoder to process facial images. By synthetically corrupting the data and forcing the neural network to map the noisy inputs back to their clean originals, the model learned a robust latent representation of human facial structures. The quantitative analysis showed an 81% reduction in Mean Squared Error, and the visual outputs confirmed the model's ability to isolate and remove random Gaussian noise while preserving essential facial features.